

```
import pandas as pd

train = pd.read_csv("train.csv")

train.head()
```

	Loan_ID	Gender	Married	Dependents	Education	Score
0	LP001002	Male	No	0	Graduate	650
1	LP001003	Male	Yes	1	Graduate	660
2	LP001005	Male	Yes	0	Graduate	670
3	LP001006	Male	Yes	0	Not Graduate	680
4	LP001008	Male	No	0	Graduate	690

Next steps:

[Generate code with train](#)

[View recommended plots](#)

```
train.shape
```

```
(614, 13)
```

```
train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               614 non-null   object
1   Gender                601 non-null   object
2   Married               611 non-null   object
3   Dependents            599 non-null   object
4   Education             614 non-null   object
5   Self_Employed         582 non-null   object
6   ApplicantIncome       614 non-null   int64
7   CoapplicantIncome     614 non-null   float64
8   LoanAmount            592 non-null   float64
9   Loan_Amount_Term      600 non-null   float64
10  Credit_History         564 non-null   float64
11  Property_Area         614 non-null   object
12  Loan_Status           614 non-null   object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
```

```
# checking for missing values
```

```
train.isna().sum()
```

ValueError

NameError X ...



Please explain the error:

NameError: name 'rf_pred' is not c



Hypothesis

The error message "NameError: name 'rf_pred' is not defined" indicates that the variable `rf_pred` has not been assigned a value before it is used in the `confusion_matrix` function. This likely occurred because you haven't trained and predicted using a model (presumably a `RandomForestClassifier`). The variable `rf_pred` is commonly used to store the predicted values from a model.

Suggested Changes



```
from sklearn.metrics import confu
from sklearn.ensemble import Ran
from sklearn.model_selection imp
```

```
# Assuming X, y are already defi
```

```
# Splitting the data into train
X_train, X_test, y_train, y_test
```

```
# Create and train the RandomFor
rf_model = RandomForestClassifie
rf_model.fit(X_train, y_train) #
```

```
# Make predictions on the test s
rf_pred = rf_model.predict(X_tes
```

```
# Now you can use rf_pred in you
print(confusion_matrix(y_test, r
print(classification_report(y_te
print("Accuracy:", accuracy_scor
```



[Use code with caution](#)

Explanation of changes:



	0
Loan_ID	Add text cell
Gender	13
Married	3
Dependents	15
Education	0
Self_Employed	32
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	22
Loan_Amount_Term	14
Credit_History	50
Property_Area	0
Loan_Status	0

df.describe()

```
train.describe() #descriptive statistics
```



	ApplicantIncome	CoapplicantIncome	LoanAmount
count	614.000000	614.000000	592.000000
mean	5403.459283	1621.245798	146.412162
std	6109.041673	2926.248369	85.587325
min	150.000000	0.000000	9.000000
25%	2877.500000	0.000000	100.000000
50%	3812.500000	1188.500000	128.000000
75%	5795.000000	2297.250000	168.000000
max	81000.000000	41667.000000	700.000000

```
train.describe(include=[object])
```



	Loan_ID	Gender	Married	Dependents	Education
count	614	601	611	599	6
unique	614	2	2	4	
top	LP002990	Male	Yes	0	Gradua
freq	1	489	398	345	4

1. Import

RandomForestClassifier: The code now imports the necessary

`RandomForestClassifier` from `sklearn.ensemble`.

2. Initialize and train the model: A

`RandomForestClassifier` object named `rf_model` is created and trained using the `fit` method with the training data (`x_train`, `y_train`).

3. Make predictions: The predict method of the trained model is used to predict the target variable for the test data (`x_test`). These predictions are stored in the `rf_pred` variable.

4. Evaluation: Now that `rf_pred` is defined, it can be used within the `confusion_matrix`, `classification_report`, and `accuracy_score` functions to evaluate the model's performance.

► Sources

Rate this answer



```
train['Loan_Status'].value_counts()
```



count Add text cell

Loan_Status	
Y	422
N	192


dtype: int64

```
import seaborn as sns
```

```
sns.countplot(x=train['Loan_Status'])
```



<Axes: xlabel='Loan_Status', ylabel='count'>



```
train['Gender'].value_counts()
```




count

Gender	
Male	489
Female	112

dtype: int64


```
sns.countplot(x=train['Gender'])
```

↗ <Axes: xlabel='Gender', ylabel='count'>




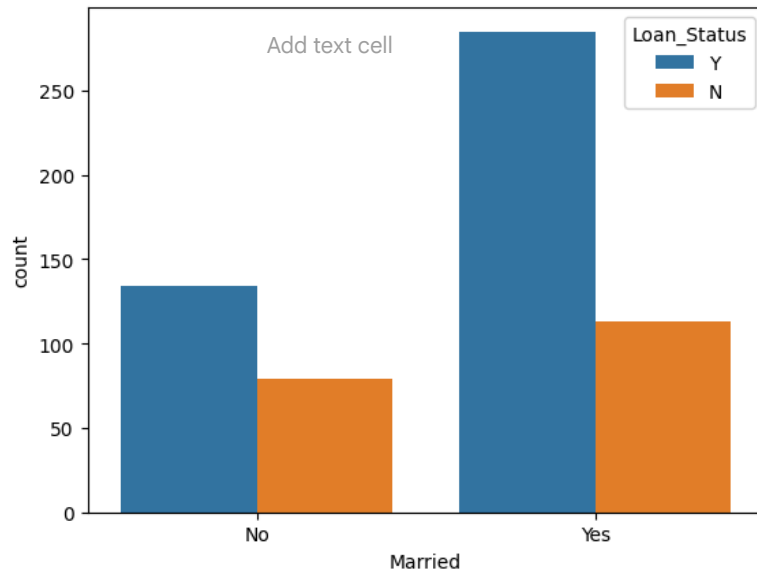
```
sns.countplot(x=train['Gender'], hue=train['Loan_Status'])
```

↗ <Axes: xlabel='Gender', ylabel='count'>




```
sns.countplot(x='Married', data=train, hue='Loan_Status')
```

 <Axes: xlabel='Married', ylabel='count'>



```
sns.displot(train['LoanAmount'], kde=True)
```

 <seaborn.axisgrid.FacetGrid at 0x7d6e9ffc24d0>



```
import matplotlib.pyplot as plt
```

```
grid = sns.FacetGrid(train, row='Gender', col='Married', he
```

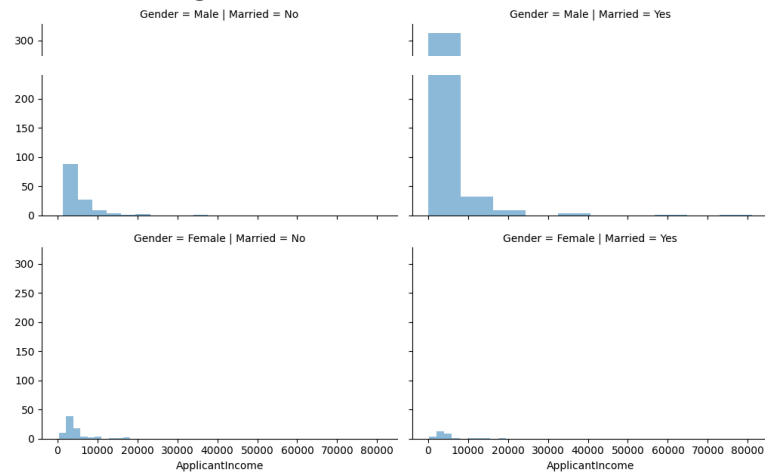
```
grid.map(plt.hist, 'ApplicantIncome', alpha=.5, bins=10)
```

Add text cell

```
grid.add_legend()
```



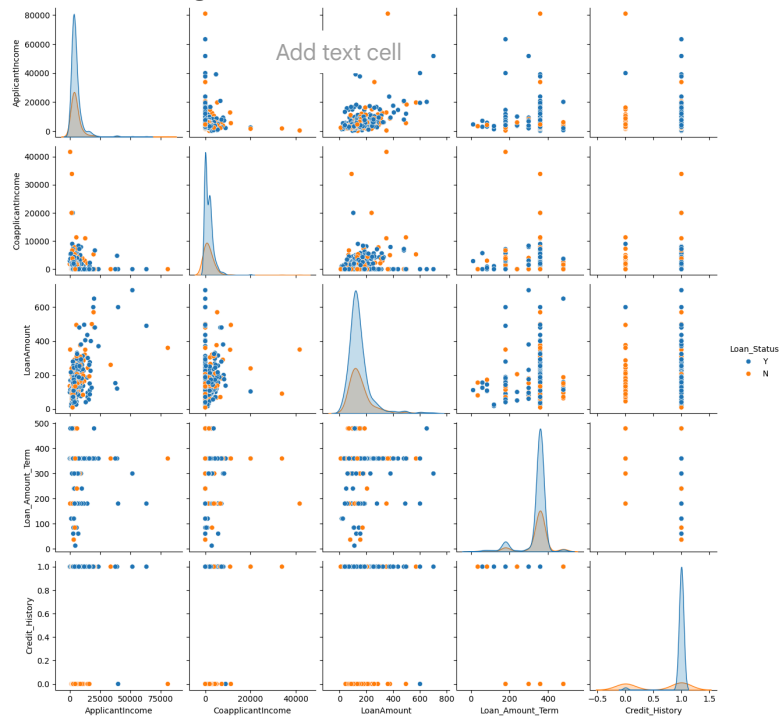
<seaborn.axisgrid.FacetGrid at 0x7d6e65798550>



```
sns.pairplot(train, hue='Loan_Status', height=2.5)
```



<seaborn.axisgrid.PairGrid at 0x7d6e63692050>



```
train.isna().sum()
```



0

Loan_ID	Add text cell
Gender	13
Married	3
Dependents	15
Education	0
Self_Employed	32
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	22
Loan_Amount_Term	14
Credit_History	50
Property_Area	0
Loan_Status	0

df['Loan_Status']

```
train['Married'] = train['Married'].fillna(train['Married']
```

```
train['Dependents'].fillna(train['Dependents'].mode()[0], i
```



<ipython-input-18-54017d2d16d3>:1: FutureWarning: A val
The behavior will change in pandas 3.0. This inplace me

For example, when doing 'df[col].method(value, inplace=

```
train['Dependents'].fillna(train['Dependents'].mode()
```

```
train['Self_Employed'].fillna(train['Self_Employed'].mode()
```



<ipython-input-19-afd59976b687>:1: FutureWarning: A val
The behavior will change in pandas 3.0. This inplace me


For example, when doing 'df[col].method(value, inplace=

```
train['Self_Employed'].fillna(train['Self_Employed'].
```

```
train['LoanAmount'].fillna(train['LoanAmount'].median(), ir
```

```
train['Loan_Amount_Term'].fillna(train['Loan_Amount_Term'].
```

```
train['Credit_History'].skew()
```


 <ipython-input-20-c03efc4edc79>:1: FutureWarning: A value is being passed to `fillna` that is not a scalar (or a list/tuple of scalars). The behavior will change in pandas 3.0. This inplace method is deprecated. For example, when doing `df[col].method(value, inplace=`


Add text cell

```
train['LoanAmount'].fillna(train['LoanAmount'].median())
<ipython-input-20-c03efc4edc79>:3: FutureWarning: A value is being passed to fillna that is not a scalar (or a list/tuple of scalars). The behavior will change in pandas 3.0. This inplace method is deprecated. For example, when doing df[col].method(value, inplace=
```

```
train['Loan_Amount_Term'].fillna(train['Loan_Amount_Term'].astype(np.float64)(-1.8823610612186696))
```

◀ ▶

```
train.isna().sum()
```



	0
Loan_ID	0
Gender	13
Married	0
Dependents	0
Education	0
Self_Employed	0
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	0
Loan_Amount_Term	0
Credit_History	50
Property_Area	0
Loan_Status	0

dtype: int64

◀ ▶

```
from sklearn.preprocessing import LabelEncoder


feature_col = ['Gender', 'Married', 'Dependents', 'Education',

le = LabelEncoder()

for col in feature_col:


    train[col] = le.fit_transform(train[col])

train.Loan_Status = train.Loan_Status.replace({"Y": 1, "N"
```

 <ipython-input-23-92b0e0b1b9aa>:1: FutureWarning: Downc
train.Loan_Status = train.Loan_Status.replace({"Y": 1

Add text cell

```
train['total_income'] = train['ApplicantIncome'] + train['C  
train.drop(columns=['ApplicantIncome', 'CoapplicantIncome']  
train.head(3)
```



	Loan_ID	Gender	Married	Dependents	Education	S
0	LP001002	1	0	0	0	
1	LP001003	1	1	1	0	
2	LP001005	1	1	0	0	

Next
steps:

[Generate code with train](#)

[View recommended plots](#)


```
rel_feat=['Gender', 'Married', 'Dependents', 'Education', '  
LoanAmount', 'Loan_Amount_Term', 'Credit_History', '  
Loan_Status', 'train_income']
```

```
rel_feat=['Gender', 'Married', 'Dependents', 'Education', '  
LoanAmount', 'Loan_Amount_Term', 'Credit_History', '  
Loan_Status', 'total_income']
```

```
# Selecting only the relevant numerical features for correl  
numerical_features = train[rel_feat].select_dtypes(include=
```

```
rel_feat_corr = numerical_features.corr()['Loan_Status']
```

```
print(rel_feat_corr) # To display correlation values.
```

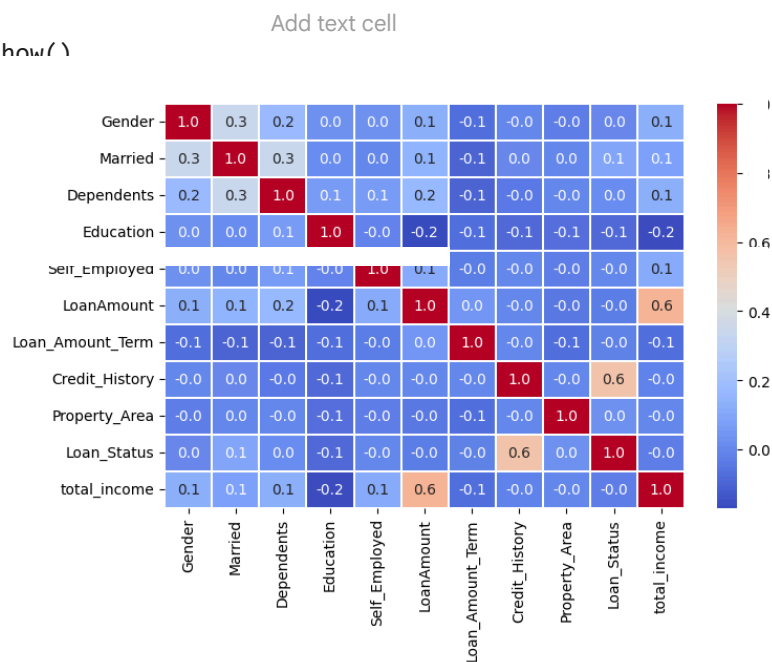


```
Gender          0.008690  
Married         0.091478  
Dependents      0.010118  
Education       -0.085884  
Self_Employed   -0.003700  
LoanAmount      -0.033214  
Loan_Amount_Term -0.022549  
Credit_History  0.561678  
Property_Area   0.032112  
Loan_Status     1.000000  
total_income    -0.031271  
Name: Loan_Status, dtype: float64
```

```
plt.figure(figsize=(8,5))
```

```
sns.heatmap(train[rel_feat].corr(), cmap='coolwarm', anno
```

nl+ show()



#Separating target variable and other variables

```
X=train.drop(columns='Loan_Status')
```

```
y=train['Loan_Status']
```

#Splitting the data into train and test sets

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(X, y, test_s
```

Enter a prompt here



0 / 2000

Responses may display inaccurate or offensive information that doesn't represent Google's views.
[Learn more](#)